

An Automated Ticket-Writing Machine!


By: Benjamin Kosiborod and Victoria Staada





Problem Statement

Given information about a car and its parking spot, we wanted to predict the type of violation that is ticketed



Dataset

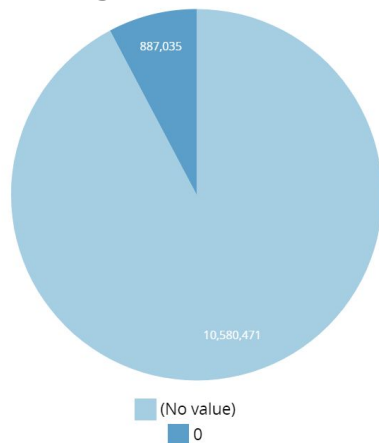
- NYC Parking Violations Issued - Fiscal Year 2019
 - ~11.5 million rows
 - 43 columns
 - 99 violation codes
 1. No Parking Street Cleaning Code 21 - **1,803,467 Violations (~16%)**
 2. Failure to Display Muni-Meter Receipt Code 38 - **1,165,883 Violations (~10%)**
 3. Photo School Zone Speed Violation Code 36 - **1,098,298 Violations (~10%)**
 4. No Standing Day Time Limits Code 14 - **1,013,584 Violations (~9%)**
 5. No Parking Day Time Limits Code 20 - **795,751 Violations (~7%)**
 - = ~51%




Approach & Methodology

- Data cleanup
 - Dropped columns we didn't need/were unhelpful (43 → 17)
 - Dropped rows with bad data for our features (~11.5M → ~4.5M)
 - Imputed clean data when applicable (e.g. standardizing abbreviations)
 - Used Open Data NYC's visualization tool to help
 - Dropped all rows with NaN values
 - Reset index
 - Clean data! 🧼
- Model training & testing
 - Sampled 100K rows from our data
 - Encoded data
 - Removed outliers
 - Trained & tested kNN, Decision Trees, Random Forest, AdaBoost, and FFNN models
 - Calculated metrics for each individual model
 - Plotted an ROC curve for all models

Unregistered Vehicles?




Vehicles Make




Challenges

- Poorly documented data
 - Days Parking In Effect (BBBBBBB, YYYYYYY, YBBYBBB, etc.)
 - 1000 vehicle colors recorded (W, WH, WHITE, WT, WHI, WHIE, etc.)
 - 12 colors
 - Red, Orange, Yellow, Green, Blue, Purple, Beige, Brown, Gray, Silver, White, Black
- Size of dataset/computing power
 - Personal computers ❌
 - Google Colab ❌
 - NEU Discovery ✓
- NEU Discovery
 - Time limits
 - Memory limits
 - Custom Environment

Days Parking In Effect



Vehicle Color



k-nearest neighbors



Default (k=5) worked best

Expectedly low accuracy... our worst model by far.

Training

	Precision	Recall	F1 Score	Support
accuracy			0.40	65661
macro avg	0.21	0.16	0.12	65661
weighted avg	0.41	0.40	0.38	65661

Testing

	Precision	Recall	F1 Score	Support
accuracy			0.17	21827
macro avg	0.04	0.04	0.04	21827
weighted avg	0.16	0.17	0.16	21827

Decision Trees



max_depth = 13 worked best for our data

Our second best model. Likely overfitted the least.

Training

	Precision	Recall	F1 Score	Support
accuracy			0.51	65661
macro avg	0.59	0.25	0.29	65661
weighted avg	0.53	0.51	0.47	65661

Testing

	Precision	Recall	F1 Score	Support
accuracy			0.41	21827
macro avg	0.21	0.13	0.15	21827
weighted avg	0.37	0.41	0.37	21827

Random Forest (bagging)



Experimented with some hyperparameters (max_depth, criterion, n_estimators)

Our best model, but likely overfits to the training data. Perhaps additional parameter tuning is in order...

Training

	Precision	Recall	F1 Score	Support
accuracy			1.00	65661
macro avg	1.00	1.00	1.00	65661
weighted avg	1.00	1.00	1.00	65661

Testing

	Precision	Recall	F1 Score	Support
accuracy			0.47	21827
macro avg	0.31	0.21	0.22	21827
weighted avg	0.43	0.47	0.42	21827

AdaBoost



Used our Decision Tree estimator from earlier as our base estimator

Unexpectedly low accuracy... Perhaps the dataset is more suited to bagging methods

Training

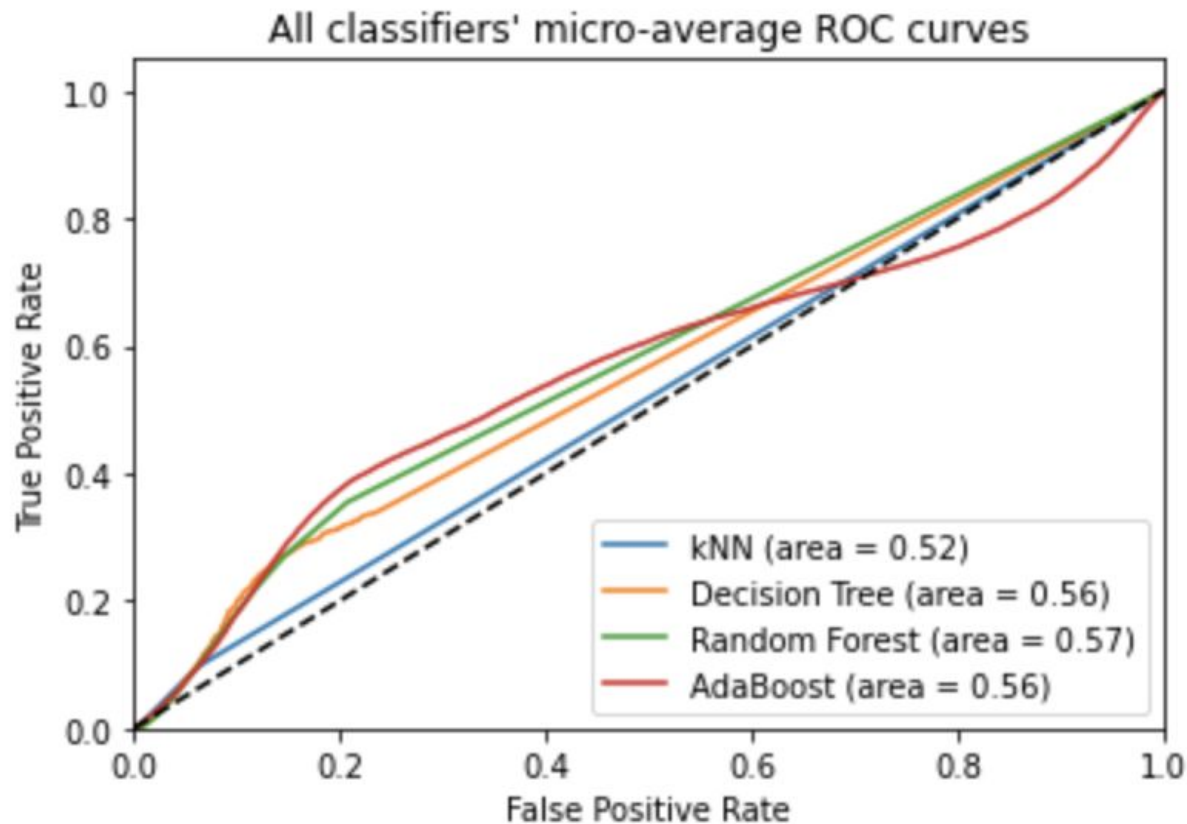
	Precision	Recall	F1 Score	Support
accuracy			0.60	65661
macro avg	0.89	0.74	0.80	65661
weighted avg	0.62	0.60	0.60	65661

Testing

	Precision	Recall	F1 Score	Support
accuracy			0.29	21827
macro avg	0.20	0.11	0.13	21827
weighted avg	0.29	0.29	0.28	21827



ROC Curve



Training duration: 57.86256742477417

1379/1379 [=====] - 2s 1ms/step - loss: 2.6884 - accuracy: 0.1932

Network's test loss and accuracy duration: 59.784929037094116



Neural Network (FFNN)

Accuracy: ~19%


Dimensions

Input size: 17 (# of features)

Output size: 99 (# of classes)

params: 159,099

Activation functions: sigmoid, softmax



Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 500)	9000
activation_24 (Activation)	(None, 500)	0
dense_25 (Dense)	(None, 250)	125250
activation_25 (Activation)	(None, 250)	0
dense_26 (Dense)	(None, 99)	24849
activation_26 (Activation)	(None, 99)	0
Total params: 159,099		
Trainable params: 159,099		
Non-trainable params: 0		

Questions?